



UNAM

# Transforming Data into Property Graphs with GraFlo

Create Property Graphs in 5 minutes with GraFlo


UNAM, Mexico  
20 November 2025

<https://tinyurl.com/graflo>

Alexander Belikov, GrowGraph  
[alexander-belikov.github.io](http://alexander-belikov.github.io)

# End-of-Workshop Achievements

1. install graflo
2. ingest a graflo example
3. create a new working schema
4. identify a limitation graflo
5. find a graflo bug
6. contribute to graflo
7. ★★★



# Plan

0. Setup
1. Knowledge Graphs: What, Why and How
2. GraFlo : User Perspective
3. GraFlo Internals
4. NB: Graph Database Zoo
5. Practicum

# Glossary

## **Labeled Property Graph (LPG):**

a data model to represent (typed) nodes, relationships, and their properties

## **Graph Database:**

an engine to store and query LPGs

## **Vertex (or Node):**

A fundamental unit of data representing an entity (e.g., a **Person**, a **Publication**). Has a type and properties.

## **Edge (or Relationship):**

A directed connection between two vertices (e.g., **owns**, **citedBy**). May have properties.

## **Graph Schema:**

A blueprint defining the allowed types of nodes, relationships, and their properties in the graph.

## **ETL (Extract, Transform, Load)**

The process of moving data from sources to a target system. GraFlo simplifies the Transformation and Loading for graphs.

## **Declarative Ingestion**

Specifying what the graph should be (via schema) rather than writing step-by-step code for how to build it.

## **Community Detection**

the process of grouping nodes into densely connected subgroups within a network

# LPG: What, Why, How

Vertices have types (labels)

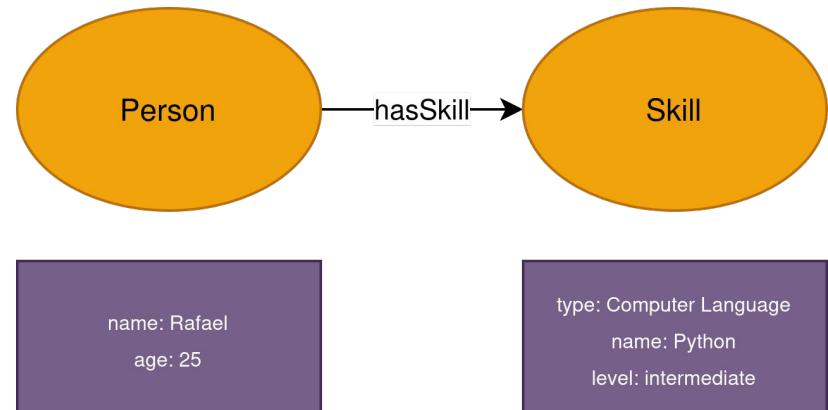
Directed relationships have types

Both have properties

Uniqueness constraints (for edges)

Data Structure + Scaffolding

Scaffolding: Query Language, Data Science tools



# LPG: What, Why, How

Graphs can be represented as Tables... But

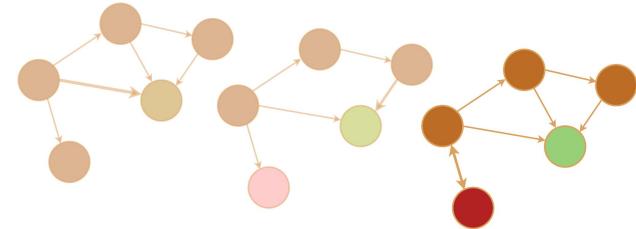
1.  Relationships: Implicit in SQL (consider 3NF), Explicit in PKG.
2.  Schema fluidity: SQL schemas are structured and require migrations for change in a PKG, these concepts are native.
3.  Superior Performance for Connections, Paths, and Patterns: recursive JOINs for SQL, that degrade exponentially, natural for Graph DBs.
4.  Holistic, Un-siloed Views of Data
5. Native AI/ML Readiness

SQL: Not designed for graph-native machine learning or structural reasoning.

PKG: The graph topology is the feature. PKGs directly enable:

- Graph Neural Networks (GNNs)
- Community and cluster detection
- Centrality and influence analysis
- Knowledge-graph grounding for LLMs
- Multi-hop reasoning and retrieval

# Typical Graph ML tasks



World Model, Digital Twin

## Node-Level Tasks

- Node Classification (Node Labeling):** Predict labels or categories for nodes (e.g., sector classification, fraud detection)
- Node Regression:** a continuous value for each node (e.g., employee churn probability, asset risk score)
- Node Embedding:** dense vector representations that capture graph context

## Edge-Level Tasks

- Link Prediction (Edge Prediction):** Predict missing or future edges (e.g., latent relationships between analysts & firms)
- Edge Classification / Regression:** Predict type/weight of relationship (e.g., sentiment strength, transaction volume)

## Graph-Level Tasks

- Graph Classification:** Predict a label for the entire graph (e.g., classify firm networks by risk profile)
- Graph Regression:** Predict a continuous value (e.g., fund performance based on its holdings graph)

## Subgraph Tasks

- Subgraph Matching / Querying:** Identify occurrences of a given pattern (e.g., insider trading-like motifs)
- Subgraph Embedding / Evaluation:** Assess local structure (e.g., influence zones, team dynamics)

## Unsupervised & Generative Tasks

- Community Detection:** Discover tightly connected groups (e.g., analyst echo chambers)
- Graph Generation:** Model how graphs evolve or simulate new structures (e.g., synthetic networks for stress tests)

# LPG: What, Why, How

## 1. Native Graph Data Structures

- **Node & Edge Storage:** Instead of rows and columns, data is stored as nodes (vertices) and relationships (edges), both capable of holding key-value properties.
- **Pre-Wired Connections:** Relationships are stored as direct pointers or links between nodes. This allows for **constant-time traversal** - jumping from one node to its neighbors is incredibly fast, unlike costly SQL JOINs.
- **Example:** A `Customer` node is directly linked to an `Order` node via a `:PURCHASED` relationship.

## 2. Scaffolding

- **Indexes:** Accelerate finding the search space for traversals. **Constraints:** Ensure data integrity and uniqueness (e.g., "Ensure `ProductID` is unique for all `Product` nodes"). This prevents duplicate data and maintains a clean graph.

## 3. Intuitive Graph Query Language

- **Declarative Pattern Matching:** You describe the **shape of the sub-graph** you're looking for, and the engine finds it for you. **Purpose:** Makes complex relationship queries intuitive and readable, turning what would be a multi-paragraph SQL query into a few clear lines of code.

# Motivational Problem: Eigenfactor calculation

## Goal:

Measure a journal's *true* scientific influence, not just how many papers it publishes.

## Idea:

A journal is influential if it is cited by **other influential journals**.

Influence “flows” through the citation network - like reputation spreading.

## How it's computed (conceptually):

1. Look at **all citations between journals** over a time window.
2. Count how often *Journal A* cites *Journal B*.  
This forms a **journal-to-journal influence matrix** (a network).
3. Repeatedly pass influence along citation link (similar to PageRank).
4. The stable flow gives each journal an **Eigenfactor score**.

## Key point:

The calculation starts by counting **all 3-hop paths**:  $\text{Journal A} \leftarrow \text{Pub} \xrightarrow{\text{-cites-}} \text{Pub} \rightarrow \text{Journal B}$ .

# Why Graph Databases Make Eigenfactor Easy

## The problem:

To build the influence matrix, we must follow millions of “who-cites-whom” links across several steps.

With files or SQL databases, this involves many complex JOINs and huge intermediate tables.

## Why it's painful in SQL or CSV files:

- Citations form a **network**, not a table.
- Multi-step (“A cites B which belongs to C...”) queries become slow and hard to write.
- Adding time filters or more hops makes queries balloon in complexity.
- Storing papers, journals, and citations in separate tables hides the structure we need.

## Why it's easy in a Graph Database:

- Citations are stored as **edges**, not JOINs.
- Walking “Paper → Cites → Paper → Journal” is a **natural graph traversal**.

## AQL query

```
FOR j IN media FILTER j.issn in ["2049-3630", ...]
RETURN MERGE({{ja : j.issn}},
{{stats : (
  FOR p in 1 INBOUND j publications_media_edges
  FILTER p.year == _year
  FOR p2 in 1 OUTBOUND p publications_publications_edges
  FILTER p2.year < _year AND p2.year >= (_year - delta)
  FOR j2 in 1 OUTBOUND p2 publications_media_edges
  FILTER j2.issn in ["2049-3630", ...]
  COLLECT jbt = j2.issn WITH COUNT INTO size SORT size DESC
  RETURN {{jb : jbt, s : size}})}})
```

# How to Ingest Data to Graph Databases?

## Tell the GDB about

- the correspondences about the fields in the data and properties of the vertices and edges
- which fields should be treated as unique
- which pairs of sets of fields form an edge between two vertices

## Pain point

**Manual ETL Mapping Overhead:** Every source requires custom scripts or queries to map raw data to nodes/edges. This becomes unmanageable as sources grow.

**Schema Drift and Rigidity:** Adapting graph schema when data models change (e.g. new fields or entity types) is error-prone and lacks abstraction.

**Duplication of Logic Across Pipelines:** Repeating transformation logic in multiple ingestion scripts leads to code duplication, versioning issues, and maintenance debt.

**No Separation of Concerns:** Data mapping, transformation, and loading are tightly coupled, making reuse or schema redesign difficult.

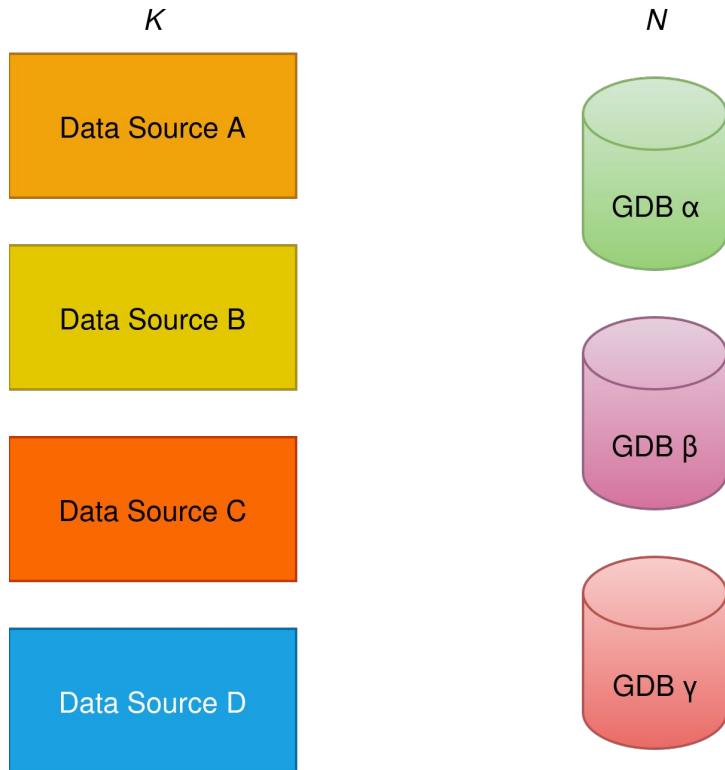
**Lack of Reusability and Documentation:** Imperative ingestion pipelines (e.g. in GSQL, Cypher, Python) are often undocumented and hard to share or audit.

**Hard to Support Multiple Backends:** Each target DB (Neo4j, TigerGraph, etc.) requires different ingestion syntax and connectors—no portability across engines.

**Poor Scalability Without Parallelism:** Ad hoc scripts often process sequentially and can't leverage multiple cores for high-volume data ingestion.

**No Declarative Validation or Indexing:** Constraints like unique IDs or compound indexes must be manually enforced or scripted, increasing risk of inconsistencies.

# How to Ingest Data to Graph Databases?




A complex web of  $K \times N$  pipelines

Clearly there must be a Graph Schema & Transformation Language (**GSTL**) that encodes separately

- transformations for each of  $K$  sources
- graph schema
- load functions for each of  $N$  databases

# Solution: Graflo



Developer Perspective: implement  $N$  backends.

User Perspective: implement  $K$  data source adapters.

## Features:

- **Declarative transformations** vs. custom ETL coding
- Adapters for Neo4j, ArangoDB and TigerGraph: Multi-database adapter eliminates vendor lock-in
- Tested on graphs with billions of edges

# GraFlo: User Perspective

## Workflow

- Study your dataset
- Create a mental model
- Create a schema
- Visualize the schema
- Set up the backend
- Ingest
- Check the ingestion results

## Schema setup

1. Define Vertices
  - a. Properties (optional types)
  - b. Indexes
2. Define Edges
  - a. Properties/Weights
  - b. Indexes
3. Define Resources
  - a. Transformations
  - b. Vertex Mappings
  - c. Edge Mappings (optional)
4. Define a Vocabulary of Transforms
5. Schema metadata

# Plotting

## plot\_schema

This command creates multiple visualizations of the schema:

1. Vertex-to- vertex relationships
2. Vertex fields and their relationships
3. Resource mappings


The visualizations are saved to the specified output path.

Args:

schema\_path: Path to the schema configuration file

figure\_output\_path: Path where the visualization will be saved

prune\_low\_degree\_nodes: Whether to remove nodes with low connectivity from the visualization (default: False)



Example: `$ uv run plot_schema -c schema.yaml -o schema.png`

Options:

-c, --schema-path PATH [required]

-o, --figure-output-path PATH

[required]

-p, --prune-low-degree-nodes BOOLEAN

# Schema: Vertex Config and Edge Config

**VertexConfig:** Defines vertex collections (nodes) in the graph

**EdgeConfig:** Defines edge collections (relationships) between vertices

**VertexConfig**

```
vertices: list[Vertex],  
blank_vertices: list[str],  
db_flavor: DBFlavor  
)
```

**EdgeConfig(edges: list[Edge])**

```
vertex_config:  
  vertices:  
    - name: person  
      fields: [id, name, age]  
      indexes: [{fields: [id]}]  
  
  edge_config:  
    edges:  
      - source: person  
        target: department
```

Extended Vertex Example

```
vertex_config:  
  vertices:  
    - name: person  
      fields: [id, name, age]  
      indexes:  
        - fields: [id]  
        - unique: false  
          fields: [name, age]  
      filters: []  
      dbname: person
```

Extended Edge Example

```
edge_config:  
  edges:  
    - source: person  
      target: department  
      relation_field: role  
      weights:  
        direct: [date, score]  
      vertices:  
        - vertex: person  
          fields: [age]  
      indexes:  
        - fields: [date]
```

# Transforms

## Setting Up Transforms in Schema

### 1. Global Transform Library (Schema Level)

Define reusable transforms in the schema's transforms section

### 2. Inline Transforms (Resource Level)


```
transforms:  
  keep_suffix_id:  
    foo: split_keep_part  
    module: graflo.util.transform  
    params:  
      sep: "/"  
      keep: -1  
    input: [id]  
    output: [_key]
```

Transforms map and transform fields during ingestion.

1. Declarative mapping: rename/remap fields (no code)
2. Functional transforms: custom Python functions for complex transformations
3. Vertex mappings (special case, check example 1)


# Example: Analyst Reports [IBES]

| TICKER | CUSIP    | CNAME          | OFTIC | ACTDATS  | ESTIMID  | ANALYST       | ERECCD | ETEXT      | IRECCD | ITEXT |
|--------|----------|----------------|-------|----------|----------|---------------|--------|------------|--------|-------|
| 0000   | 87482X10 | TALMER BANCORP | TLMR  | 20140310 | RBCDOMIN | ARFSTROM J    | 2      | OUTPERFORM | 2      | BUY   |
| 0000   | 87482X10 | TALMER BANCORP | TLMR  | 20140311 | JPMORGAN | ALEXOPOULOS S | NaN    | OVERWEIGHT | 2      | BUY   |
| 0000   | 87482X10 | TALMER BANCORP | TLMR  | 20140311 | KEEFE    | MCGRATTY C    | 2      | OUTPERFORM | 2      | BUY   |



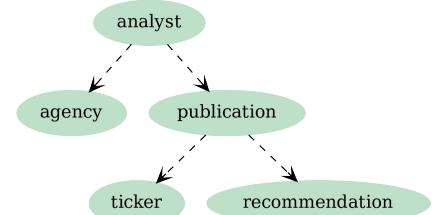
Institutional Brokers' Estimate System (IBES)

# IBES Schema: Vertices



```
ibes.yaml x : ibes.yaml x
43 vertex_config:
44     blank_vertices:
45         - publication
46     vertices:
47         - name: publication
48             dbname: publications
49             fields:
50                 - datetime_review
51                 - datetime_announce
52             indexes:
53                 - fields:
54                     - _key
55                 - type: hash
56                     unique: false
57                     fields:
58                         - datetime_review
59                 - type: hash
60                     unique: false
61                     fields:
62                         - datetime_announce
63
64 vertex_config:
65     vertices:
66         - name: ticker
67             dbname: tickers
68             fields:
69                 - cusip
70                 - cname
71                 - oftic
72             indexes:
73                 - fields:
74                     - cusip
75                     - cname
76                     - oftic
77             name: agency
78             dbname: agencies
79             fields:
80                 - aname
81             indexes:
82                 - fields:
83                     - aname
```

# IBES Schema: Vertices & Edges



```
43   vertex_config:          103   edge_config:
44     vertices:              104     edges:
45       - name: analyst      105       - source: publication
46         dbname: analysts  106       target: ticker
47         fields:           107       - source: analyst
48           - last_name    108       target: agency
49           - initial      109       weights:
50         indexes:          110       vertices:
51           - fields:        111       - name: publication
52             - last_name   112       keep_vertex_name: false
53             - initial    113       fields:
54           - name:        114       - datetime_review
55             dbname:      115       - datetime_announce
56             recommendations 116       - source: analyst
57             fields:        117       target: publication
58               - erec        118       - source: publication
59               - etext      119       target: recommendation
60               - irec        120
61               - itext      121
62             indexes:        122
63               - fields:      123
64                 - irec      124
65                 - erec      125
66                 - etext      126
67                 - itext      127
68   edge_config:            128
```

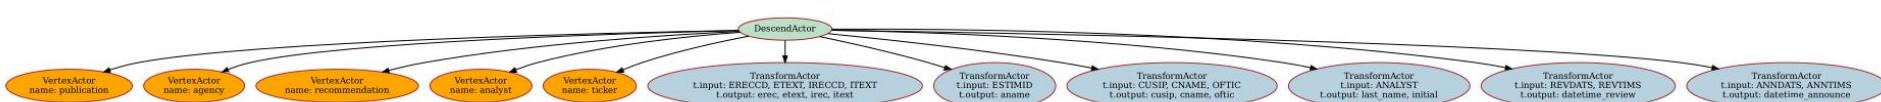
# IBES Schema: Transforms

```

3   resources:                                     ✓ 21 ^ v
4     - resource_name: ibes
5       apply:
6         - vertex: ticker
7         - vertex: analyst
8         - vertex: recommendation
9         - vertex: agency
10        - foo: parse_date_ibes
11          module: graflo.util.transform
12          input:
13            - ANNDATS
14            - ANNTIMS
15          output:
16            - datetime_announce
17        - foo: parse_date_ibes
18          module: graflo.util.transform
19          input:
20            - REVDATS
21            - REVTIMS
22          output:
23            - datetime_review
24

3   resources:
4     - resource_name: ibes
5       apply:
6         - foo: cast_ibes_analyst
7           module: graflo.util.transform
8           input:
9             - ANALYST
10            output:
11              - last_name
12              - initial
13            map:
14              CUSIP: cusip
15              CNAME: cname
16              OFTIC: oftic
17            map:
18              ESTIMID: aname
19            map:
20              ERECCD: erec
21              ETEXT: etext
22              IRECCD: irec
23              ITEXT: itext
24

```



# Extra Functions

Beyond ingestion, GraFlo provides utilities for querying, filtering, and managing graph data.

**fetch\_present\_documents()** - Check which documents from a batch already exist

**keep\_absent\_documents()** — Get documents that don't exist in the database

**aggregate()** — Perform aggregations on collections

- Supports COUNT, SUM, AVG, MIN, MAX
- Group by a discriminant field
- Apply filters

# GraFlo internals: Resource Transformation to Vertices & Edges


Actors are processing units that transform raw data into graph structures.

Raw data is a json (potentially nested).

They run in sequence within a Resource, each handling a specific transformation step.

Descend actor is responsible for recursive processing and mapping.

Four actor types, executed in priority order:



# GraFlo internals: Resource Transformation to Vertices & Edges

Actors are processing units that transform raw data into graph structures. They run in sequence within a Resource, each handling a specific transformation step.

**Actor Types:** four actor types, executed in priority order

**DescendActor** (priority 10):

- Processes nested/hierarchical data structures
- Expands lists or nested dictionaries
- Executes child actors for each nested item
- Maintains location tracking via LocationIndex

**TransformActor** (priority 20):

- Applies data transformations
- Executes transform functions (e.g., field mapping, data cleaning)
- Results stored in buffer\_transforms or buffer\_vertex
- Can target specific vertices

**VertexActor** (priority 50): Creates vertex documents

- Extracts fields from documents
- Merges data from transforms and buffers
- Accumulates vertices in ctx.acc\_vertex[vertex\_name][lindex]
- Applies filters and field selection

**EdgeActor** (priority 90): Creates edges between vertices

- Merges vertices first (via merge\_vertices())
- Renders edges based on vertex matches
- Calculates edge weights from vertex fields

# Graph Databases

| Database   | License    | Query Language | Horizontal Scaling    | Comment                 |
|------------|------------|----------------|-----------------------|-------------------------|
| JanusGraph | Apache 2.0 | Gremlin        | Native                | Best for massive scales |
| Nebula     | Apache 2.0 | nGQL           | Native                |                         |
| Dgraph     | Apache 2.0 | GraphQL        | Native                |                         |
| Apache AGE | Apache 2.0 | Cyber+SQL      | PostgreSQL clustering |                         |
| Memgraph   | BSL        | Cypher         | Enterprise only       | In-memory focus         |
| Neo4j      | GPLv3      | Cypher         | Enterprise only       | Mature Ecosystem        |
| ArangoDB   | BSL        | AQL            | Enterprise only       | Close to noSQL          |
| TigerGraph | Prop       | GSQL           | Enterprise only       | Typed Schema            |

# Graph DB Extras

## Web Interface

- **Neo4j:** Bloom (visual exploration), Browser (Cypher queries)
- **ArangoDB:** Web UI (queries, management, graph visualization)
- **TigerGraph:** GraphStudio (visual schema design, query building)

## Graph Data Science

- **Neo4j: GDS Library** - 70+ production-grade algorithms (PageRank, community detection, node similarity, machine learning pipelines)
- **ArangoDB: Graph Analytics** - Custom implementations using Pregel, integrated with machine learning connectors
- **TigerGraph: Built-in GSQl** - 30+ parallel graph algorithms optimized for massive-scale analytics

## Key Differentiators

- **Neo4j:** Most mature GDS ecosystem, enterprise support
- **ArangoDB:** Multi-model analytics (document + graph)
- **TigerGraph:** Native parallel execution on distributed data

# Graph DB Idiosyncrasy

## Neo4j

Pure Labeled Property Graph. Vertices have labels (types). Indexes over multiple properties. Types not enforced but possible to set constraints. Type mismatches may only surface at runtime.  
Schema-last.

## ArangoDB

Document-Graph Hybrid models. Types are collections (of JSONs). Edges store handles to nodes.  
Automatic index on `_key` (can be manipulated). Can represent nested data.  
Type agnostic but possible to enable type validation.

## Tigergraph

Schema-first, rigid, strongly-typed. Secondary indexes can be only single property. Graphs have to be composed from globally defined vertices and edges.

## Nomenclature.

# Application: Reviewer Recommendation

## Authors

```
author_id,FullName,HIndex,research_sector  
309238221625,Guillaume Lemaître,10,32057259  
747324850364,Patrick L. Meras,4,8258574  
987843024183,S. I. Konovalov,5,30262949
```

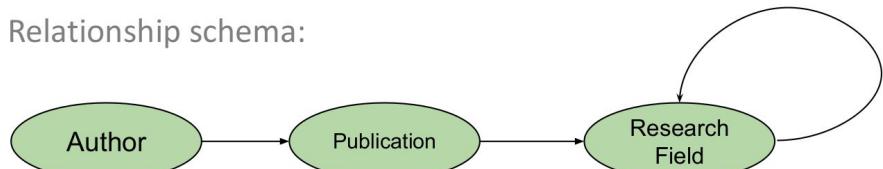
## Research Fields

```
field_of_study_id,display_name,level,children  
87687168,Digital audio,4,6800068  
87687168,Digital audio,4,24579023  
87687168,Digital audio,4,30246029
```

## Publications

```
PublicationId,authors,topics,publication_year,Doi  
465031,"['id:300648343950, name:Tadeusz Kaczorowski'  
'id:566936217113, name:Anna-Karina Kaczorowska'  
'id:214748948560, name:Sebastian Dorawa']", [185592680 89423630], 2019, 10.3390/V11070657
```

Relationship schema:




# Application: Reviewer Recommendation

## Problems

1. The same author may have multiple ids.
2. Given the publication records, who would be a good reviewer for paper?

## Solutions

1. Disambiguate
  - a. construct **[coAuth] Author ↔ Author** from Author → Publication
  - b. derive communities
  - c. fuzzy match within communities
2. Recommend Reviewers
  - a. construct **[coDomain] Author ↔ Author** from Author → Publication → ResearchField
  - b. derive communities in **[coDomain] Author ↔ Author**
  - c. For each new publication pick reviewers from the same **[coDomain]** comm id that have a different **[coAuth]** comm id



# Practicum

Check your .env files (ports specifically, but also username/password)

Arangodb Web Interface: <http://localhost:8535>

Neo4j Web Interface: <http://localhost:7475>

Try to ingest examples from <https://github.com/growgraph/graflo/tree/main/examples>  
and visualize them using Web Interface.

# Practice Creating Schemas

1. [grouplens.org/datasets/movielens/20m/](http://grouplens.org/datasets/movielens/20m/)
2. <https://www.kaggle.com/datasets/rmisra/news-category-dataset>
3. <https://www.kaggle.com/datasets/mylesoneill/magic-the-gathering-cards>
4. [secret dataset]

# Roadmap

1. Improve API UX (how easy is it for developer to use the package).
2. Implement SQL schema to GraFlo schema generator
3. Add SQL API: ingestion of SQL resources
4. Add Schema validation and version control
5. Add Nebula and Janus as Graph Database backends.

# Conclusion

- The World is a Graph: The most complex and connected data is naturally modeled as a graph.
- GraFlo Makes it Practical: It eliminates the ETL bottleneck, providing a declarative, scalable, and multi-database framework for building your knowledge graph.
- You Are Now Equipped: You can install GraFlo, design schemas, ingest data, and leverage the power of graph databases in minutes, not days.